

AI506 Term Project – Final Report

Sangjune Kim
KAIST
Daejeon, South Korea

Gyuchan An
KAIST
Daejeon, South Korea

Minseo Kang
KAIST
Daejeon, South Korea

ABSTRACT

We study advertisement click behavior in a real-world search-log dataset through two prediction tasks. Task 1 (click prediction) is an extremely imbalanced binary classification problem (click rate $\approx 1.1\%$) whose validation set is dominated by cold users; Task 2 (ad prediction) is a zero-history retrieval problem ranking the clicked ad among 17,518 candidates. For Task 1 we show that the direct text similarity between a query and a displayed ad barely separates clicks from non-clicks, and we instead exploit a collaborative signal we call *Semantic K-Nearest-Neighbor Click Prediction* (SKNCP): we look up the ads that were clicked by the most similar past queries and measure how similar the current ad is to them. We empirically validate the smoothness assumption that this signal relies on (similar queries click on similar ads), and use SKNCP plus trusted click-through-rate (CTR) features as a small correction to a click-tuned content score, reaching a clicked-class F1 of 0.1109 on the external validation set (bootstrap mean 0.104, 95% CI [0.067, 0.144]), exceeding the provided baseline (0.0436) by +154%. For Task 2 we represent queries, ads, and users in a shared text embedding space and maintain multiple per-user interest vectors updated online; this reaches NDCG@3 = 0.1538, improving over a text-similarity-only baseline by +55%. We choose deterministic Task 1 components (weights, K , and threshold) on an internal split of the training data and report the final external score using the frozen content-score cache included with the submission files.

1 INTRODUCTION

Online advertising platforms select advertisements based on search queries, user context, and ad metadata, and a user’s click is a sparse but direct signal of interest. In this project we study this setting through two related prediction tasks defined over the same real-world search-log dataset. The two tasks are formally distinct and require independent submissions; accordingly our work proceeds along two parallel tracks with task-specific models motivated by separate exploratory findings.

1.1 Problem Definition

Common notation. Let $\mathcal{U}$, $\mathcal{S}$, and $\mathcal{A}$ denote the sets of users, search events, and advertisements. For a search event $s \in \mathcal{S}$, let $u_s \in \mathcal{U}$ be the issuing user, $x_s \in \mathbb{R}^d$ its *query embedding*, and for an ad $a \in \mathcal{A}$ let $z_a \in \mathbb{R}^d$ be its *ad-title embedding*, with $d = 384$. Embeddings are precomputed with a pretrained sentence encoder (MiniLM) and are *frozen*. The training log is a set of displayed-ad impressions $\{(s_i, u_{s_i}, a_i, r_i, h_i, y_i)\}_{i=1}^N$ with $N = 320,000$, where r_i is the display position, h_i the engine-provided historical CTR, and $y_i \in \{0, 1\}$ the click label. For a search event s , we write a_s^+ for its clicked ad (when one exists) and $z_{a_s^+}$ for its embedding; this lets us distinguish a user’s *search intent* x_s from the *clicked-ad representation* $z_{a_s^+}$, a distinction the Task 2 model relies on.

Terms used throughout.

- **CTR** (click-through rate): fraction of impressions that are clicked. The *global* training CTR is $\mu \approx 0.0111$.
- **HistCTR** $h_{s,a}$: the historical CTR estimate *supplied by the search engine* for the displayed (ad, position); an input feature.
- **Smoothed entity CTR**: for an entity key g (an AdID, IPID, device, or category), the train-only Laplace-smoothed click rate $\text{ctr}(g) = (\text{clicks}(g) + k\mu) / (\text{count}(g) + k)$ with $k = 20$; unseen keys fall back to μ .
- **Internal split / external validation.** Model choices (features, hyper-parameters, decision thresholds) are made on an *internal* split of the 320k training log: we sort by SearchID and use the first 80% (256,157 rows) as internal-train and the last 20% (63,843 rows) as internal-validation. The *external validation* set is the separately provided `click_validation` file (20,000 impressions); it is held out for final F1/AUC measurement.
- **Evaluation.** Task 1 uses clicked-class $F_1 = 2 \text{ TP} / (2 \text{ TP} + \text{FP} + \text{FN})$; Task 2 uses NDCG@3. *AUC* (area under the ROC curve) measures ranking quality independently of threshold. *Cohen’s d* denotes a standardized mean difference (effect size).
- Task-specific terms are defined formally where first used, in Sections 2 and 3.

Task 1: Click Prediction.

- **GIVEN:** a displayed-ad impression (s, a) with search-event ID, ad ID, position $r_{s,a}$, and HistCTR $h_{s,a}$.
- **FIND:** a binary prediction $\hat{y}_{s,a} \in \{0, 1\}$ of whether a is clicked in s .
- **TO MAXIMIZE:** clicked-class F_1 .

Task 2: Ad Prediction.

- **GIVEN:** a search event s whose clicked ad is hidden, plus user, search-context, ad metadata, and the query/ad embeddings.
- **FIND:** three distinct ads ranked by likelihood of being the clicked ad.
- **TO MAXIMIZE:** NDCG@3.

The candidates are the full pool of 17,518 ads.

1.2 Exploratory Data Analysis Findings

Table 1 collects the statistics that drive our design.

Table 1: Exploratory statistics motivating the design choices.

Quantity	Value
Training impressions / clicks / CTR	320,000 / 3,560 / 1.11%
External validation impressions / clicks / CTR	20,000 / 229 / 1.15%
Train $\cap$ external SearchID overlap	1 ($\approx 0.01\%$)
Cold users in click-prediction validation	$\approx 29\%$
Users / users with zero clicks	14,133 / 83.1%
Ads with zero observed clicks	87.1%
Click rate at position 1 vs. 7	1.29% / 0.71%
Non-logged-in vs. logged-in click rate	1.20% / 0.99%
Ads per search (1 / 2)	57% / 43%
Mean HistCTR, clicked / non-clicked	0.022 / 0.0095
Task 1 new users / ads in val.-test	34.4–45.8% / 9.0–11.7%
Search–ad category match rate	53.1%
CTR on category match / mismatch	0.86% / 1.39%
Task 1 direct query–ad cosine (clicked vs. non)	0.527 / 0.500 ($d \approx 0.11$)
Task 2 query–(clicked vs. random) ad cosine	0.55 / 0.25 ($d \approx 1.56$)

Three properties matter most. (i) *Clicks are extremely rare* (1.11%), so clicked-class F1 is dominated by where the rare positives land at the top of the ranking, not by overall accuracy; moreover, 83.1% of users and 87.1% of ads have zero observed clicks, limiting pure collaborative filtering. (ii) *The validation set is a cold, new-session regime*: external SearchIDs essentially never appear in training (1 overlap), $\approx 29\%$ of validation users are unseen, and the validation/test sets introduce many new users and ads (34.4–45.8% and 9.0–11.7%, respectively), so session-memorizing or user-specific features do not transfer. (iii) *Text similarity behaves very differently in the two tasks*, and this distinction is central to our Task 1 design. In Task 2 the relevant comparison is *clicked ad vs. a random ad* for the same query: clicked ads are much more query-similar than random ones (mean cosine 0.55 vs. 0.25, Cohen’s $d \approx 1.56$), so embedding similarity is a strong retrieval signal. In Task 1 the relevant comparison is *clicked vs. non-clicked among the already-displayed ads*: here the direct query–ad cosine barely separates the two classes (0.527 vs. 0.500, Cohen’s $d \approx 0.11$), because the displayed ads are already topically plausible. Task 1 therefore needs a signal beyond direct text similarity, which motivates the collaborative SKNCP signal in Section 2.

Two additional tabular priors are useful for Task 1. First, HistCTR is the strongest supplied per-impression signal: clicked impressions have mean HistCTR 0.022, compared with 0.0095 for non-clicked ones. Second, category equality is not a reliable relevance rule: search and ad categories match in 53.1% of training impressions, but the click rate is lower on matching pairs (0.86%) than on non-matching pairs (1.39%), partly because many search contexts have an unspecified category. We therefore use smoothed category interactions and trusted CTR priors rather than a naive match indicator.

These observations yield two qualitatively different problems. Task 1 is a class-imbalanced classification problem over cold entities with at most two ads per search (so within-search ranking structure is weak). Task 2 is a zero-history retrieval problem in which the

only signal that generalizes is the shared embedding space. We address them with independent task-specific models.

2 TASK 1: CLICK PREDICTION

2.1 Motivation: from failed direct similarity to SKNCP

A natural idea is to score an impression by the cosine similarity $\cos(x_s, z_a)$ between the query and ad embeddings. As noted above, this fails for Task 1: among displayed impressions the clicked and non-clicked classes are nearly indistinguishable in direct cosine (external Cohen’s $d \approx 0.15$, $AUC \approx 0.52$ – 0.54). The displayed ads are already topically matched to the query, so direct text similarity says little about which one is clicked.

We therefore reference *what similar users actually clicked*. We define the **Semantic K-Nearest-Neighbor Click Prediction** (SKNCP) signal:

$$\text{SKNCP}(s, a) = \max_{a' \in C_K(s)} \cos(z_a, z_{a'}), \quad (1)$$

where $C_K(s)$ is the set of ads that were *clicked* in the K training search events whose query embeddings are most similar to x_s . In words: “find the K past searches most like this one, collect the ads those searches clicked, and measure how similar the current ad is to the closest of them.”

Empirical justification. SKNCP relies on a *smoothness assumption*: *semantically similar queries click on similar ads*. We validate this directly. Sampling $\approx 30,000$ click pairs and binning by query–query similarity, the average similarity between their clicked ads increases monotonically with query similarity (Figure 1, Pearson $r = 0.336$); the reverse direction (binning by clicked–ad similarity and measuring query similarity) gives the same effect ($r = 0.341$). The query space and the clicked-ad space thus move together (a joint structure), which is exactly the condition a K -NN approach needs. We then verify that the resulting signal is strong: for 2,000 clicked samples, the max cosine between the current clicked ad and the clicked ads of its $K=100$ most similar queries averages 0.879, versus 0.674 for random ads (Figure 2a, Cohen’s $d = 1.27$), and a per-click paired comparison shows SKNCP exceeding the random baseline for $\approx 80\%$ of clicks (Figure 2b). On its own SKNCP is only modestly stronger than the direct cosine for click prediction (external Cohen’s $d \approx 0.24$ vs. 0.15), consistent with its weakness as a stand-alone classifier (Section 2.4). Its value is as a *feature*: combined with the trusted-CTR, position, and rank signals the log-linear decision score reaches external Cohen’s $d \approx 0.61$ —roughly $4\times$ the direct-similarity baseline.

SKNCP alone is nevertheless a *weak* stand-alone classifier (Section 2.4): its index only contains ads that were actually clicked in training ($\approx 1,759$ ads, $\approx 67\%$ coverage), so it cannot score impressions whose neighborhood is empty. We therefore use SKNCP as a collaborative-filtering *feature* that complements the sparse trusted CTR signals, not as a model by itself.

2.2 Methodology

Signals. The auxiliary log-linear component combines signals that are each justified by the EDA: (1) log HistCTR, the supplied prior; (2–5) the log of smoothed AdID, IPID, device, and category



Figure 1: Smoothness assumption (Evidence 1). Each point is a click pair: as the query–query embedding similarity grows, the similarity between the ads those queries clicked grows monotonically (binned means and linear fit, $r = 0.336$; the reverse direction gives $r = 0.341$).




Figure 2: K-NN signal validity (Evidence 2). (a) The current clicked ad is far more similar to the clicked ads of its similar queries (mean 0.879) than to random ads (mean 0.674), Cohen’s $d = 1.27$. (b) Per-click paired comparison: for $\approx 80\%$ of clicks the SKNCP similarity exceeds the random baseline.

CTRs (defined in Section 1); (6) an *inverse-propensity-score (IPS) position correction* $\log(\text{ctr}(\text{pos})/\mu)$ that removes the placement bias measured in Table 1 (position 1 vs. 7); (7) the within-search HistCTR rank; (8) SKNCP from Eq. (1); and (9) a click-tuned *content head*. The content head is the cosine of two learned linear projections of x_s and z_a , trained on clicks with a class-balanced cross-entropy and ensembled over multiple random initializations. Unlike raw query–ad cosine, this click-tuned text signal is discriminative by itself (Figure 3). We then combine the content head with trusted CTR priors into a cached *content-strong* broad score; this broad score is the “content” term used in the final blend.¹ A fixed multiplicative login term $\log(1.7)$ is added for non-logged-in rows (EDA: their CTR is higher).

¹The content head uses stochastic training, so re-training it can produce a different content-score draw and a slightly different point F1 within the bootstrap confidence interval. To make the submitted result reproducible, the reported point F1 0.1109 uses the frozen pre-computed content-score cache submitted with the report (code/task1/cs_cache.npz, MD5 3dc2105. . .).




Figure 3: Click-tuned content-head evidence on external validation. Left: clicked impressions shift to higher learned content-head scores than non-clicked impressions (AUC 0.675). Right: the click rate increases with score decile, with the top decile reaching 3.3% vs. the 1.15% base CTR.

Content-head validation. Figure 3 explains why the learned content head is retained as the broad semantic signal. On the external validation impressions, the pure content head reaches AUC 0.675, much stronger than raw query–ad cosine (≈ 0.52 – 0.54). Its highest score decile has a click rate of 3.3%, nearly three times the base click rate of 1.15%. This does not solve the top- k F1 problem by itself, because clicks are still extremely rare, but it provides the only dense signal that transfers to cold users and ads. The content-strong score used in the final blend adds the stable CTR priors to this head, yielding AUC 0.713 and F1 0.106 before the SKNCP correction.

Fitting. We weight the signals to directly maximize clicked-class F1 by *coordinate ascent*: each weight is swept over a grid while the others are held fixed, choosing the value that maximizes top- k F1 on the internal-validation split, cycling until convergence. This optimizes the metric of interest rather than log-loss, which is dominated by the abundant negatives. All smoothed CTRs and the SKNCP index are built from internal-train only when scoring internal-validation, and from the full training log when scoring external validation, so a row never sees its own label.

Final model: content-strong score with an SKNCP/CTR correction. The cached content-strong score is the most stable broad ranking signal, while SKNCP, CTR, position, and rank are sparse but useful near the top of the ranked list. We therefore do not let the F1-fitted log-linear model replace the content-strong score. Instead, we use it as a small auxiliary correction in standardized (z-scored) space:

$$\text{score}(s, a) = z(\text{content-strong}) + \alpha z(s_{LL}), \quad \alpha = 0.2, \quad (2)$$

where s_{LL} (the *log-linear score*) is the F1-fit log-linear combination of the signals above. Inside this auxiliary score, SKNCP receives the largest fitted weight (3.0), but the auxiliary score is multiplied by $\alpha=0.2$ before being added to the main content-strong score. Thus SKNCP is used as a top-of-list refinement rather than as the dominant global ranking signal. The blend weight $\alpha=0.2$ is a fixed a-priori constant; the neighborhood size $K=200$ is selected on the internal split by the procedure of Section 2.3 (not tuned on external labels). Predictions take the top fraction of scores equal to the training click rate (1.11%); this threshold uses no external labels. Figure 4 summarizes the pipeline.




Figure 4: Final Task 1 pipeline: a frozen content-strong score corrected by a small F1-optimized auxiliary score built from trusted CTR, position, rank, and SKNCP signals.

2.3 Selecting the neighborhood size K

The SKNCP feature has one hyper-parameter: the neighborhood size K , i.e. how many nearest training queries contribute clicked ads to Eq. (1). We select it on the internal split alone, and we select it by a deliberately chosen criterion: the *internal-validation Cohen’s d of the boosted decision score* (the standardized mean separation between the clicked and non-clicked classes), maximized over a grid $K \in \{50, 100, \dots, 400\}$.

We select K on Cohen’s d rather than on top- k F1 because top- k F1 is computed from integer true-positive counts among only a few hundred positives, so as a function of K it is a high-variance estimator unsuitable for hyper-parameter selection. Cohen’s d uses full-sample means and spread and gives a far more stable selection signal. The criterion is fixed before any external label is read: the external column of Table 2 is the result of the selected K , not an input to choosing it.

The internal-validation d is maximized at $K=200$ ($d=0.7342$), which is the value used by the final model. Two robustness observations accompany this choice. First, the criterion is flat: d varies by under 0.012 across the entire grid, so K is not a sensitive hyper-parameter and the selection is not sensitive. Second, every other K yields external F1 = 0.1064, within the bootstrap confidence interval of the selected operating point (Section 2.4), so the reported 0.1109 sits on a stable plateau rather than a single tuned spike. $K=200$ is chosen because it maximizes the internal d ; the external F1 of 0.1109 follows from that choice.

2.4 Results

Table 3 reports external-validation performance. The deterministic weights, threshold, and SKNCP neighborhood size are chosen on the internal split; the reported point score uses the submitted frozen content-score cache. SKNCP alone is weak (F1 0.027), confirming it is a complementary feature rather than a classifier. The

Table 2: Neighborhood-size sweep for the boosted model. Selection uses only the internal-validation Cohen’s d (leftmost numeric column); its argmax is $K=200$. External columns are reported results, not selection inputs. Produced by `models/m06_sknpc/experiments/k_sweep_eval.py`.

K	selection	reported (external)	
	int-val d	F1	AUC
50	0.7240	0.1020	0.7132
100	0.7235	0.1064	0.7131
150	0.7235	0.1064	0.7131
200	0.7342	0.1109	0.7130
250	0.7325	0.1064	0.7133
300	0.7235	0.1064	0.7131
350	0.7235	0.1064	0.7131
400	0.7235	0.1064	0.7131

Table 3: Task 1 external validation (clicked-class F1). Threshold = training click rate ($k=222$) unless noted. [†] Content-strong row uses an internal-validation positive rate as threshold ($k \approx 616$); its precision/recall are not comparable to the other rows.

Configuration	F1	Prec.	Recall	AUC
Provided baseline (HistCTR>0.01)	0.0436	—	—	—
HistCTR only	0.064	—	—	0.685
SKNCP only	0.027	—	—	0.573
Content-strong baseline [†]	0.106	0.073	0.196	0.713
SKNCP-boosted (final)	0.1109	0.113	0.109	0.713

content-strong baseline already reaches F1 0.106. The final SKNCP-boosted model attains $F_1 = 0.1109$ (AUC 0.713), a +154% relative improvement over the provided baseline (0.0436). A 2000-resample bootstrap gives a mean F1 of 0.104 with a wide 95% confidence interval [0.067, 0.144]. The interval is wide by construction: the external set contains only 229 positives, and the model predicts $k = 222$ clicks with TP = 25. A single true-positive swing corresponds to $\Delta F_1 \approx 0.004$, so the raw variance of the point estimate is high; differences below ≈ 0.04 are within sampling noise. We therefore report AUC as the more stable ranking metric alongside F1.

A stratification of the validation users by their training history, computed in the content-strong analysis before the final SKNCP correction, makes the role of the semantic component concrete. On *warm* users (with click history, 21%), the model is strongest (F1 0.12); on *cold* users (unseen, 29%), the trusted user-specific CTRs are unavailable and the click-tuned content score carries the prediction (F1 0.11). The largest and hardest stratum is search-history-only users (49%), where ranking quality is high (AUC 0.73) but positives are hard to push above the global threshold. This is the empirical reason the final model needs both a generalizing semantic signal and the SKNCP/CTR refinement.

3 TASK 2: AD PREDICTION

3.1 Motivation and Approach

The proposed method is motivated by the retrieval nature of Task 2 and the sparsity of explicit click signals in the dataset. In this task, the model must predict which advertisement was clicked for a given search event. Click history provides the most direct evidence of user preference, because a clicked advertisement indicates an explicit positive interaction. However, exploratory analysis shows that click events are sparse for many users, making it difficult to construct reliable personalized representations from clicks alone (among 14,133 users in the training set, 83.1% of users, or 11,748 users, have no click events). In contrast, search events are more broadly available and each query is represented by a pre-computed text embedding. Therefore, we use search embeddings as an additional behavioral signal to compensate for the sparsity of click history.

For each search event, the query embedding represents the semantic context of the user’s search. This makes query–advertisement semantic compatibility a natural signal for ranking candidate advertisements. However, query-level relevance alone is not sufficient for personalized advertisement prediction. The same or similar queries can be issued by different users with different interests, and these users may click different advertisements under similar search contexts. This suggests that Task 2 requires not only matching advertisements to the given query, but also modeling user-specific preference.

To address this, we represent each user with **multiple interest** vectors rather than a single global embedding. These vectors are constructed using both search query embeddings and clicked advertisement embeddings. Search queries provide dense but relatively weak behavioral evidence, while clicked advertisements provide sparse but stronger preference evidence. In the update process, clicked advertisements are assigned larger weights than search queries so that explicit positive feedback has a stronger influence on the user representation. This allows the model to build a user-specific multi-interest profile even when click history is limited.

Based on this motivation, our method ranks candidate advertisements by combining query-level relevance and user-level preference. Query–advertisement similarity captures how well a candidate advertisement matches the given search context, while advertisement–interest similarity captures how well the advertisement matches the user’s learned multi-interest profile. Because queries, advertisements, and user interests are all represented in the same embedding space, the model can perform lightweight ranking using cosine similarity without additional projection layers or complex feature engineering.

3.2 Methodology

Our method is a lightweight multi-interest retrieval model that directly uses the provided text embeddings. The main idea is to represent each user with multiple interest vectors and to rank advertisements by combining the current query context with the user’s historical preference.

For each user (u), we maintain (K) interest vectors,

$$V_u = v_{u,1}, v_{u,2}, \dots, v_{u,K}, \quad v_{u,i} \in \mathbb{R}^d.$$




Figure 5: Task 2 multi-interest model architecture.

When a new behavioral embedding ($e \in \mathbb{R}^d$) is observed for user u , the model updates all interest vectors using a soft-assignment mechanism. The embedding e can be either a search query embedding x_s or a clicked advertisement embedding $z_{a_s^+}$. We first compute the cosine similarity between e and each interest vector:

$$c_i = \cos(v_{u,i}, e).$$

The assignment weight of each interest vector is computed by a temperature-scaled softmax:

$$w_i = \frac{\exp(c_i/\tau)}{\sum_{j=1}^K \exp(c_j/\tau)},$$

where τ is the temperature parameter. A smaller τ makes the update more concentrated on the most similar interest vector, while a larger τ distributes the update more evenly.

Each interest vector is updated as

$$v_{u,i} \leftarrow \text{norm}(v_{u,i} + \eta(e)w_i\hat{e}),$$

where

$$\hat{e} = \frac{e}{\|e\|_2}$$

is the normalized event embedding, and $\text{norm}(\cdot)$ denotes L2 normalization. The update strength $\eta(e)$ depends on the type of behavioral signal:

$$\eta(e) = \begin{cases} \alpha_{\text{search}}, & \text{if } e = x_s, \\ \alpha_{\text{click}}, & \text{if } e = z_{a_s^+}. \end{cases}$$

Since search events are dense but weak signals, we use a small α_{search} . In contrast, clicked advertisements are sparse but stronger positive feedback, so we use a larger α_{click} . This allows the model to use search history to reduce click sparsity while still giving stronger influence to explicit click behavior.

For prediction, the model computes two relevance signals for a search event s and a candidate advertisement a . The first signal measures query-level semantic relevance:

$$r(s, a) = \cos(x_s, z_a).$$

The second signal measures user-level preference by comparing the candidate advertisement with the most relevant interest vector of the user u_s :

$$p(u_s, a) = \max_{1 \leq i \leq K} \cos(v_{u_s, i}, z_a).$$

The final score combines query-level relevance and user-level preference:

$$\text{score}(s, a) = (1 - \gamma_{u_s})r(s, a) + \gamma_{u_s}p(u_s, a),$$

where γ_{u_s} is an adaptive mixing weight determined by the available user history.

We define γ_{u_s} as

$$\gamma_{u_s} = \begin{cases} \gamma_{\text{click}}, & \text{if } u_s \in \mathcal{U}_{\text{click}}, \\ \gamma_{\text{search}}, & \text{if } u_s \notin \mathcal{U}_{\text{click}} \text{ and } u_s \in \mathcal{U}_{\text{search}}, \\ 0, & \text{otherwise,} \end{cases}$$

Table 4: Adaptive mixing weight used in the scoring function: The actual values used in our experiments are included.

User type	Condition	γ_{u_s}
Clicked user	$u_s \in \mathcal{U}_{\text{click}}$	0.7
Search-only user	$u_s \notin \mathcal{U}_{\text{click}}, u_s \in \mathcal{U}_{\text{search}}$	0.5
Cold-start user	otherwise	0.0

where $\mathcal{U}_{\text{click}}$ is the set of users with at least one clicked advertisement in the training data, and $\mathcal{U}_{\text{search}}$ is the set of users with search history.

If $\gamma_{u_s} = 0$, the user has no available behavioral history, and the model falls back to query-only ranking:

$$\text{score}(s, a) = r(s, a).$$

Otherwise, the user-interest score is computed by comparing the candidate advertisement with the most relevant interest vector $p(u_s, a)$ defined above. Users with click history receive the largest contribution from the learned interest profile, users with only search history receive a moderate contribution, and cold-start users are ranked only by query-advertisement similarity. For each validation search event we rank all 17,518 candidates and submit the top three:

$$\text{Top3}_{a \in \mathcal{A}} \text{score}(s, a).$$

This formulation keeps all computations in the shared embedding space of search queries, advertisements, and user interests. Therefore, the model can perform efficient ranking using cosine similarity without additional neural projection layers or complex feature engineering.

3.3 Results

Setup. We train on 320,000 (search, ad, click) triples using the given data order. Evaluation uses `ad_validation`: 214 queries, one ground-truth ad each, ranked against all 17,518 candidates.

Baselines.

- **Random:** uniform random score.
- **HistCTR:** rank by historical CTR within category (official baseline).
- **Query-only ($\gamma=0$):** pure query-ad text similarity, no user interest.

Table 5: Ad recommendation results (NDCG@3).

Method	NDCG@3	Rank-1 hits
Random	0.0008	0
HistCTR (baseline)	0.0211	2
Query-only ($\gamma=0$)	0.0989	14
MultiInterest (ours)	0.1538	17

Table 6: Rank distribution of ground-truth ads ($n=214$).

Rank	Count	NDCG weight
1	17	1.000
2	6	0.631
3	2	0.500
> 3	189	0.000

The query-only baseline already surpasses HistCTR by 4.7×, confirming that **text embedding similarity is the dominant relevance signal** for this task.

To analyze the contribution of each component, we conduct an ablation study on Task 2 using NDCG@3 as the evaluation metric. The full model uses $K = 5$, $\alpha_{\text{search}} = 0.01$, $\alpha_{\text{click}} = 0.5$, $\tau = 0.1$, $\gamma_{\text{click}} = 0.7$, and $\gamma_{\text{search}} = 0.5$.

Table 7: Ablation study on Task 2.

Model variant	NDCG@3	Difference
Full model	0.1538	–
w/o interest	0.0989	-0.0550
w/o tiered γ	0.1351	-0.0187
w/o search signal	0.1415	-0.0123
w/o click signal	0.0866	-0.0673

The ablation results show that user interest modeling is important for Task 2. Removing the interest component reduces NDCG@3 from 0.1538 to 0.0989, which corresponds to the query-only setting. This confirms that the learned user-interest vectors provide additional information beyond query-advertisement similarity.

Among the behavioral signals, the click signal has the largest contribution. When the click update is removed, NDCG@3 drops to 0.0866, which is even lower than the query-only baseline. This suggests that search-only interest vectors can introduce noisy personalization when they are not anchored by explicit click feedback. The search signal also contributes to performance, but its effect is smaller. Removing the search update decreases NDCG@3 from 0.1538 to 0.1415, indicating that search history is useful as a dense auxiliary signal but is weaker than explicit click feedback.

The tiered adaptive mixing weight also improves performance. Without γ_{search} , the model cannot effectively use search-only users, and NDCG@3 decreases from 0.1538 to 0.1351. This result supports our design choice of assigning different mixing weights to clicked users, search-only users, and cold-start users.

We also evaluate the effect of the number of interest vectors (K).

Table 8: Effect of the number of interest vectors.

K	NDCG@3	Difference
1	0.1392	-0.0146
3	0.1392	-0.0146
5	0.1538	–
10	0.1362	-0.0176

The best performance is obtained when ($K=5$). Smaller values such as ($K=1$) and ($K=3$) are less effective because they provide limited capacity to represent multiple user interests. However, increasing (K) to 10 also reduces performance, likely because click feedback is sparse and too many interest vectors make specialization more difficult.

4 CONCLUSION

We presented two methodologically independent, EDA-aligned models for the same advertising search log. For Task 1, the key insight is that direct query–ad text similarity is uninformative among displayed impressions (external Cohen’s $d \approx 0.15$); SKNCP alone is only modestly stronger ($d \approx 0.24$), but combined with CTR, position, and rank signals it lifts the decision score to $d \approx 0.61$ —about $4\times$ the direct-similarity baseline; we validated the underlying smoothness assumption directly and used SKNCP with trusted CTR and position-bias-corrected signals as a small correction to a click-tuned content score, reaching a clicked-class F1 of 0.1109 (bootstrap mean 0.104, 95% CI [0.067, 0.144]; +154% over the provided baseline), with deterministic weights, K , and threshold chosen on an internal split and the frozen content-score cache included with the submission. For Task 2, representing queries, ads, and users in a shared text embedding space and maintaining online multi-interest vectors reaches NDCG@3 = 0.1538, $7.3\times$ improvement over a HistCTR baseline. Across both tasks the shared embedding space is the signal that generalizes to the cold, new-session evaluation regime, but the two tasks exploit it differently: Task 2 uses it directly for retrieval, while Task 1 must route it through the collaborative SKNCP signal because direct similarity does not separate clicks among already-displayed ads. The dominant residual limitation is the scarcity of positive labels (229 on the click-prediction validation set), which makes small F1 differences statistically indistinguishable and motivates reporting AUC alongside F1.

REFERENCES

- [1] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu. LightGBM: A Highly Efficient Gradient Boosting Decision Tree. In *NeurIPS*, 2017.
- [2] H. B. McMahan, G. Holt, D. Sculley, M. Young, D. Ebner, J. Grady, L. Nie, T. Phillips, E. Davydov, D. Golovin, S. Chikkerur, D. Liu, M. Wattenberg, A. M. Hrafnkelsson, T. Boulos, and J. Kubica. Ad Click Prediction: A View from the Trenches. In *KDD*, 2013.
- [3] N. Reimers and I. Gurevych. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *EMNLP-IJCNLP*, 2019.

A APPENDIX

A.1 Labor Division

The team performed the following tasks for this final report.

- Dataset audit and consolidation of EDA findings [Sangjune Kim, Gyuchan An, Minseo Kang].
- Task 1 model design, feature engineering, and validation [Gyuchan An].
- Task 2 multi-interest model design, streaming update logic, and validation [Sangjune Kim].
- Report integration, template formatting, and consistency checks [Minseo Kang].

A.2 Full Disclosure

Sangjune Kim. He is not doing any dissertation, thesis, or external project related to this class project.

Gyuchan An. He is not doing any dissertation, thesis, or external project related to this class project.

Minseo Kang. He is not doing any dissertation, thesis, or external project related to this class project.

Use of LLM assistance. Large language models, including ChatGPT and Claude, were used for LaTeX editing, wording refinement, and organization of experimental findings. The experiments, numerical results, modeling decisions, and final submissions were produced and verified by the team.